

logistic regression

For this activity, you work as a consultant for an airline. The airline is interested in knowing if a better in-flight entertainment experience leads to higher customer satisfaction. They would like you to construct and evaluate a model that predicts whether a future customer would be satisfied with their services given previous customer feedback about their flight experience

```
In [1]: 1
        2 # Standard operational package imports.
        3 import numpy as np
        4 import pandas as pd
        5
        6 # Important imports for preprocessing, modeling, and evaluation.
        7 from sklearn.preprocessing import OneHotEncoder
        8 from sklearn.model_selection import train_test_split
        9 from sklearn.linear_model import LogisticRegression
       10 import sklearn.metrics as metrics
       11
       12 # Visualization package imports.
       13 import matplotlib.pyplot as plt
       14 import seaborn as sns
```

```
In [2]: 1 # Load dataset
        2
        3 df = pd.read_csv("air.csv")
```

In [3]:

1df.head(10)

Out[3]:

	satisfaction	Customer Type	Age	Type of Travel	Class	Flight Distance	Seat comfort	Departure/Arrival time convenient	Food and drink	log Likelihood
0	satisfied	Loyal Customer	65	Personal Travel	Eco	265	0	0	0	
1	satisfied	Loyal Customer	47	Personal Travel	Business	2464	0	0	0	
2	satisfied	Loyal Customer	15	Personal Travel	Eco	2138	0	0	0	
3	satisfied	Loyal Customer	60	Personal Travel	Eco	623	0	0	0	
4	satisfied	Loyal Customer	70	Personal Travel	Eco	354	0	0	0	
5	satisfied	Loyal Customer	30	Personal Travel	Eco	1894	0	0	0	
6	satisfied	Loyal Customer	66	Personal Travel	Eco	227	0	0	0	
7	satisfied	Loyal Customer	10	Personal Travel	Eco	1812	0	0	0	
8	satisfied	Loyal Customer	56	Personal Travel	Business	73	0	0	0	
9	satisfied	Loyal Customer	22	Personal Travel	Eco	1556	0	0	0	

10 rows × 22 columns

Explore the data

Check the data type of each column. Note that logistic regression models expect numeric data.

```
In [4]: 1 df.dtypes
```

```
Out[4]: satisfaction                object
Customer Type                     object
Age                               int64
Type of Travel                    object
Class                            object
Flight Distance                   int64
Seat comfort                      int64
Departure/Arrival time convenient int64
Food and drink                   int64
Gate location                    int64
Inflight wifi service            int64
Inflight entertainment           int64
Online support                   int64
Ease of Online booking           int64
On-board service                 int64
Leg room service                 int64
Baggage handling                 int64
Checkin service                  int64
Cleanliness                      int64
Online boarding                  int64
Departure Delay in Minutes       int64
Arrival Delay in Minutes         float64
dtype: object
```

To predict customer satisfaction, check how many customers in the dataset are satisfied before modeling.

```
In [5]: 1 df['satisfaction'].value_counts(dropna = False)
```

```
Out[5]: satisfied      71087
dissatisfied    58793
Name: satisfaction, dtype: int64
```

There were 71,087 satisfied customers and 58,793 dissatisfied customers.

54.7 percent (71,087/129,880) of customers were satisfied. While this is a simple calculation, this value can be compared to a logistic regression model's accuracy.

An assumption of logistic regression models is that there are no missing values. Check for missing values in the rows of the data.

```
In [6]: 1 df.isnull().sum()
```

```
Out[6]: satisfaction                0
Customer Type                      0
Age                                0
Type of Travel                     0
Class                              0
Flight Distance                    0
Seat comfort                        0
Departure/Arrival time convenient  0
Food and drink                     0
Gate location                      0
Inflight wifi service              0
Inflight entertainment             0
Online support                     0
Ease of Online booking             0
On-board service                   0
Leg room service                   0
Baggage handling                   0
Checkin service                    0
Cleanliness                        0
Online boarding                    0
Departure Delay in Minutes         0
Arrival Delay in Minutes           393
dtype: int64
```

For this activity, the airline is specifically interested in knowing if a better in-flight entertainment experience leads to higher customer satisfaction. The Arrival Delay in Minutes column won't be included in the binomial logistic regression model; however, the airline might become interested in this column in the future.

```
In [7]: 1 # Drop the rows with missing values and save the resulting pandas DataFrame
2
3
4 df_subset = df.dropna(axis=0).reset_index(drop = True)
```

If you want to create a plot (sns.regplot) of your model to visualize results later in the notebook, the independent variable Inflight entertainment cannot be "of type int" and the dependent variable satisfaction cannot be "of type object."

Make the Inflight entertainment column "of type float."

```
In [8]: 1
2 df_subset = df_subset.astype({"Inflight entertainment": float})
```

Convert the categorical column satisfaction into numeric through one-hot encoding.

```
In [9]: 1 OneHotEncoder(drop='first').fit_transform(df_subset[['satisfaction']])
```

```
Out[9]: <129487x1 sparse matrix of type '<class 'numpy.float64'>'
        with 70882 stored elements in Compressed Sparse Row format>
```

```
In [10]: 1 OneHotEncoder(drop='first').fit_transform(df_subset[['satisfaction']]).toa
```

```
Out[10]: array([[1.],
                [1.],
                [1.],
                ...,
                [0.],
                [0.],
                [0.]])
```

```
In [11]: 1 OneHotEncoder().fit_transform(df_subset[['satisfaction']]).toarray()
```

```
Out[11]: array([[0., 1.],
                [0., 1.],
                [0., 1.],
                ...,
                [1., 0.],
                [1., 0.],
                [1., 0.]])
```

```
In [12]: 1
          2 df_subset['satisfaction'] = OneHotEncoder(drop='first').fit_transform(df_s
```

```
In [13]: 1 # To examine what one-hot encoding did to the DataFrame, output the first
          2
          3
          4 df_subset.head(10)
```

Out[13]:

	satisfaction	Customer Type	Age	Type of Travel	Class	Flight Distance	Seat comfort	Departure/Arrival time convenient	Food and drink	Is
0	1.0	Loyal Customer	65	Personal Travel	Eco	265	0	0	0	
1	1.0	Loyal Customer	47	Personal Travel	Business	2464	0	0	0	
2	1.0	Loyal Customer	15	Personal Travel	Eco	2138	0	0	0	
3	1.0	Loyal Customer	60	Personal Travel	Eco	623	0	0	0	
4	1.0	Loyal Customer	70	Personal Travel	Eco	354	0	0	0	
5	1.0	Loyal Customer	30	Personal Travel	Eco	1894	0	0	0	
6	1.0	Loyal Customer	66	Personal Travel	Eco	227	0	0	0	
7	1.0	Loyal Customer	10	Personal Travel	Eco	1812	0	0	0	
8	1.0	Loyal Customer	56	Personal Travel	Business	73	0	0	0	
9	1.0	Loyal Customer	22	Personal Travel	Eco	1556	0	0	0	

10 rows × 22 columns



```
In [14]: 1 # you ca see the result from the satisfaction column
```

Create the training and testing data

Put 70% of the data into a training set and the remaining 30% into a testing set. Create an X and y DataFrame with only the necessary variables.

```
In [15]: 1 X = df_subset[["Inflight entertainment"]]
          2 y = df_subset["satisfaction"]
          3
          4 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.3, ra
```

```
In [16]: 1 X.head()
```

```
Out[16]:
```

	Inflight entertainment
0	4.0
1	2.0
2	0.0
3	4.0
4	3.0

Model building

```
In [17]: 1 clf = LogisticRegression().fit(X_train,y_train)
```

Obtain parameter estimates Make sure you output the two parameters from your model.

```
In [18]: 1 clf.coef_
```

```
Out[18]: array([[0.99751462]])
```

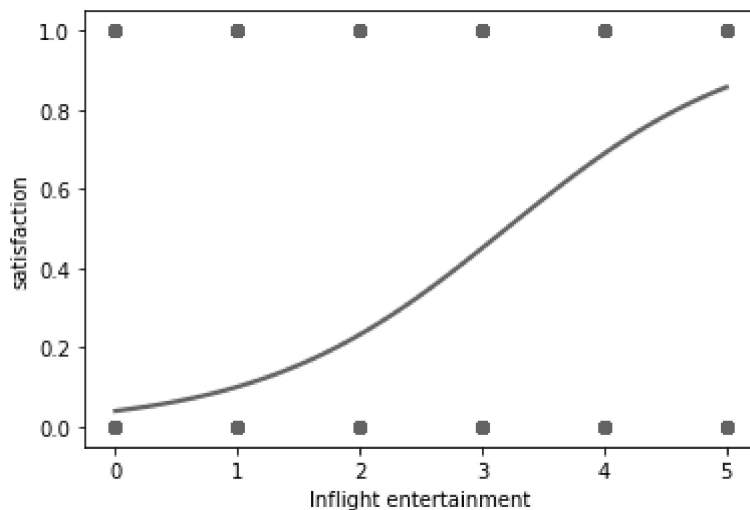
```
In [19]: 1 clf.intercept_
```

```
Out[19]: array([-3.19355406])
```

Create a plot of your model to visualize results using the seaborn package.

```
In [20]: 1 sns.regplot(x="Inflight entertainment", y="satisfaction", data=df_subset,
```

```
Out[20]: <AxesSubplot:xlabel='Inflight entertainment', ylabel='satisfaction'>
```



The graph seems to indicate that the higher the inflight entertainment value, the higher the customer satisfaction, though this is currently not the most informative plot. The graph currently doesn't provide much insight into the data points, as Inflight entertainment is categorical.

Results and evaluation

Predict the outcome for the test dataset

```
In [25]: 1 # Save predictions.  
2  
3 y_pred = clf.predict(X_test)
```

```
In [29]: 1 # In order to examine the predictions, print out y_pred.  
2  
3  
4 print(y_pred)
```

```
[1. 0. 0. ... 0. 0. 0.]
```

```
In [30]: 1 # Use predict_proba to output a probability.  
2 clf.predict_proba(X_test)
```

```
Out[30]: array([[0.14258068, 0.85741932],  
                [0.55008402, 0.44991598],  
                [0.89989329, 0.10010671],  
                ...,  
                [0.89989329, 0.10010671],  
                [0.76826225, 0.23173775],  
                [0.55008402, 0.44991598]])
```

```
In [31]: 1 # Use predict to output 0's and 1's.  
2  
3  
4 clf.predict(X_test)
```

```
Out[31]: array([1., 0., 0., ..., 0., 0., 0.])
```

Analyze the results


```
In [32]: 1  ### YOUR CODE HERE ###
2
3  print("Accuracy:", "%.6f" % metrics.accuracy_score(y_test, y_pred))
4  print("Precision:", "%.6f" % metrics.precision_score(y_test, y_pred))
5  print("Recall:", "%.6f" % metrics.recall_score(y_test, y_pred))
6  print("F1 Score:", "%.6f" % metrics.f1_score(y_test, y_pred))
```

Accuracy: 0.801529
 Precision: 0.816142
 Recall: 0.821530
 F1 Score: 0.818827

Precision measures the proportion of data points predicted as True that are actually True.
 (Ranges from 0 to 1)


Recall measures the proportion of data points that are predicted as True, out of all the data points that are actually True.(Ranges from 0 to 1)

Accuracy measures the proportion of data points that are correctly classified.(Ranges from 0 to 1)

The model performed well since "Accuracy", "Precision", "Recall" and "F1 Score" are up to 80 %

```
In [35]: 1  # let's examine the type of error made by the algorithm
2
3  cm = metrics.confusion_matrix(y_test, y_pred, labels = clf.classes_)
4  disp = metrics.ConfusionMatrixDisplay(confusion_matrix = cm,display_labels
5  disp.plot())
```

Out[35]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1c2cf594be0>



1.True negatives: 13714

The count of observations that a classifier correctly predicted as False (0)

2. True positives: 17423

The count of observations that a classifier correctly predicted as True (1)

3. False positives: 3925

The count of observations that a classifier incorrectly predicted as True (1)

4. False negatives: 3785

The count of observations that a classifier incorrectly predicted as False (0)

True positive - False negative = $17423 - 3925 = 13498$

True negative - False negative = $13714 - 3785 = 9929$

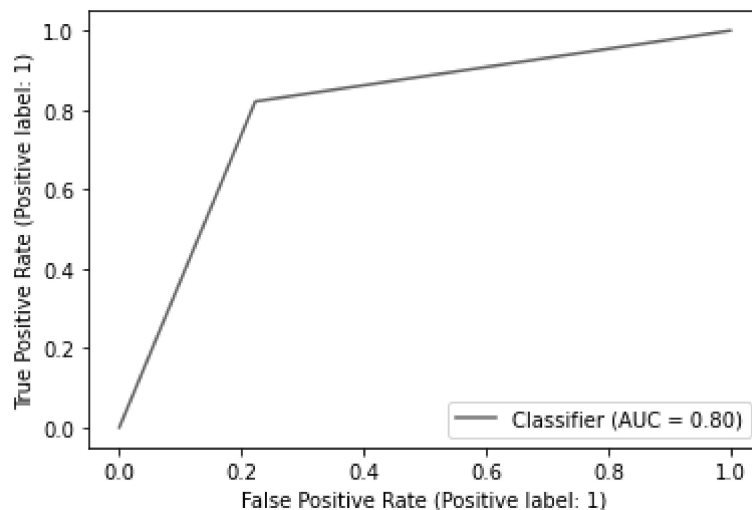
The model performance is not too bad

ROC curves

An ROC curve helps in visualizing the performance of a logistic regression classifier.

```
In [41]: 1 from sklearn.metrics import RocCurveDisplay
```

```
In [42]: 1 RocCurveDisplay.from_predictions(y_test, y_pred)
2 plt.show()
```



In this graph, the ROC curve indicates that the corresponding classifier performs decently well.
AUC = 80 %

Findings and Recommendations

Customers who rated in-flight entertainment highly were more likely to be satisfied. Improving in-flight entertainment should lead to better customer satisfaction.

The model is 80.2 percent accurate. This is an improvement over the dataset's customer satisfaction rate of 54.7 percent.

The success of the model suggests that the airline should invest more in model development to examine if adding more independent variables leads to better results. Building this model could not only be useful in predicting whether or not a customer would be satisfied but also lead to a

In []:

1	
---	--